

# Salient Object Detection (SOD) Project Report

**Author:** Ridon Maliqi

**Date:** November 2025

## Project Report – Salient Object Detection (SOD)

### 1. Project Overview

- Goal: Develop a **Salient Object Detection system from scratch** using CNNs.
- Saliency focuses on identifying the **most visually important objects** in an image.
- Implemented the **full pipeline**: dataset preprocessing, model design, training, evaluation, and interactive demo.

### 2. Learning Objectives

By completing this project, I achieved:

- Strengthened understanding of **deep learning fundamentals**.
- Implemented a **complete ML pipeline** without pre-trained models.
- Handled image datasets efficiently, including **augmentation**.
- Designed, trained, and debugged a **CNN encoder-decoder architecture**.
- Computed **IoU, Precision, Recall, F1-score, and MAE**.
- Visualized model predictions and interpreted results critically.
- Built a **Gradio app** for interactive demonstration.

### 3. Technical Requirements

#### 3.1 Dataset

- Dataset: **DUTS** (22000+ images, pixel-accurate saliency masks)
- Preprocessing:
  - Resize to 224×224
  - Normalize pixels to 0–1
  - Train/Validation/Test split: 70% / 15% / 15%
  - Augmentations: horizontal flip, random crop, brightness adjustment

#### 3.2 Model Architecture

- Input: RGB image (224×224×3)
- Encoder: 4 Conv2D layers with ReLU, each followed by MaxPooling
- Decoder: 4 ConvTranspose2D layers with ReLU
- Output: 1-channel Sigmoid mask
- Loss: Binary Cross-Entropy +  $0.5 \times (1 - \text{IoU})$
- Optimizer: Adam ( $\text{lr} = 1\text{e-}3$ )
- Epochs: 20, Early stopping applied

#### 3.3 Training

- Custom **training loop** with:
  - Forward/backward pass
  - Logging of loss and metrics per epoch
  - Validation per epoch
  - Checkpoint saving and resuming capability

#### 3.4 Evaluation

- Metrics: IoU, Precision, Recall, F1-score, MAE
- Visualizations: Input → Ground Truth → Predicted Mask → Overlay

#### 3.5 Experiments & Improvements

- Improvements over baseline:
  - Added **Batch Normalization** and **Dropout**
  - Increased depth of convolutional layers

- Enhanced augmentation (horizontal flip, random crop, brightness shift)
- Compared **baseline vs improved** in table:

Model	IoU	Precision	Recall	F1
Baseline	0.6504	0.7831	0.8484	0.8144
Improved	0.6449	0.8238	0.8049	0.8143

### 3.6 Demo / Visualization

- **Gradio web app** implemented:
  - Upload an image → display predicted saliency mask → overlayed visualization
  - Measures **inference time per image (~0.025 sec)**
  - User-friendly interface for interactive demonstration

## 4. Tools & Environment

- Language: Python 3.9+
- Framework: TensorFlow/Keras
- Libraries: NumPy, OpenCV, Matplotlib, scikit-learn, tqdm, Gradio
- Environment: Google Colab (T4 GPU), Local Jupyter Notebook for testing
- Version Control: Git/GitHub

## 5. Experiments and Results

- Training progress monitored via **loss and IoU curves**
- Sample visual results:

### Example Visualizations:

1. Input Image
  2. Ground-truth Mask
  3. Predicted Mask
  4. Overlay
- Observed that the **baseline model** performed better than the improved version in IoU; training further is needed for improvements to surpass baseline

## 6. Learnings

- Building SOD from scratch enhances **deep learning intuition**
- Proper **dataset preprocessing and augmentation** is critical
- Visualization helps **debug model predictions**
- Gradio provides **easy deployment** for interactive demos
- Early stopping and checkpoints ensure **training stability and reproducibility**

## 7. Deliverables

- **Code files:** `data_loader.py`, `sod_model.py`, `train.py`, `evaluate.py`, `app.py`
- **Report:** 8 pages PDF
- **Presentation Slides:** 5 slides summarizing key results